

EECS 391: Introduction to AI

Soumya Ray

Website: http://engr.case.edu/ray_soumya/eecs391_sp15/

Email: sray@case.edu

Office: Olin 516

Office hours: F 11:30-2pm or by appointment

Announcements

- HW due next Thursday

Today

- Goal-Directed Search (Chapter 3)

Environment Type

- We'll assume the environment is:
 - Fully observable (to track the state)
 - Static (shouldn't change while agent is searching)
 - Deterministic (agent needs to be able to precisely predict states resulting after each action)
- Some search algorithms also need discrete environments

Problem Setup

- Our agent is currently in some state of the world
 - Call this the *initial state*
- It wants to get to a different state of the world
 - Call this the *goal state*
 - In general, the desired target may be defined by a logical predicate (*goal test*) that encompasses a *set* of goal states
 - In this case the agent wants to get to *any* goal state satisfying the goal test

Problem Setup (2)

- To change the state of the world, the agent has actions
 - These will be called “search operators”
 - Also called “successor functions”, same thing
- The search operators applied successively to the initial state generate a sequence of states
 - This is the “search space”

Problem Setup (3)

- Each search operator has an associated **cost**
- The **search objective** is to discover a sequence of operators that takes the agent from the initial state to any goal state with minimum cost

Checklist

- Initial state
- Goal Test
- Search Operator (agent “action”/successor fn)
- Operator/Step Cost
- Objective: Find a sequence of actions that get from initial state to goal with minimum cost

Basic Steps of Search Algorithms

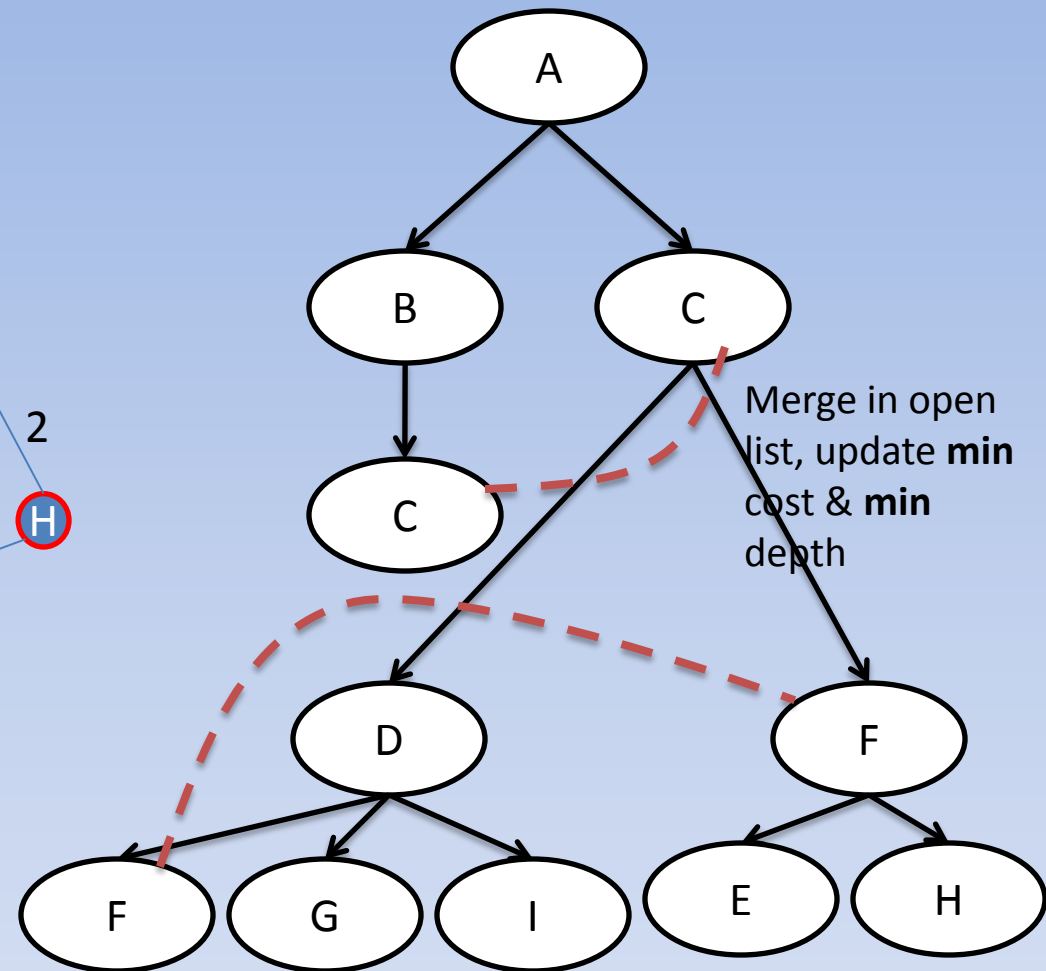

- Add initial state to **open list** (list of unvisited states)
- While open list is not empty
 - **Remove node from the open list**
 - If this is the goal, solution found, return path
 - Else
 - Find the successors of this node (“**expanding a state**”)
 - Add the current state to the list of visited (expanded) states (**closed list**)
 - Add the new states to the open list (if they do not appear in the closed list)
 - State could already be in open list, set parent pointer correctly (based on minimum path cost)

Search algorithms differ in order of node removal

Blind Search 1: Breadth First Search

- Shallowest (lowest depth) nodes on open list are expanded first
- First expand successors of initial state, then their successors, etc
- Usually, the open list will be implemented via a queue

Example



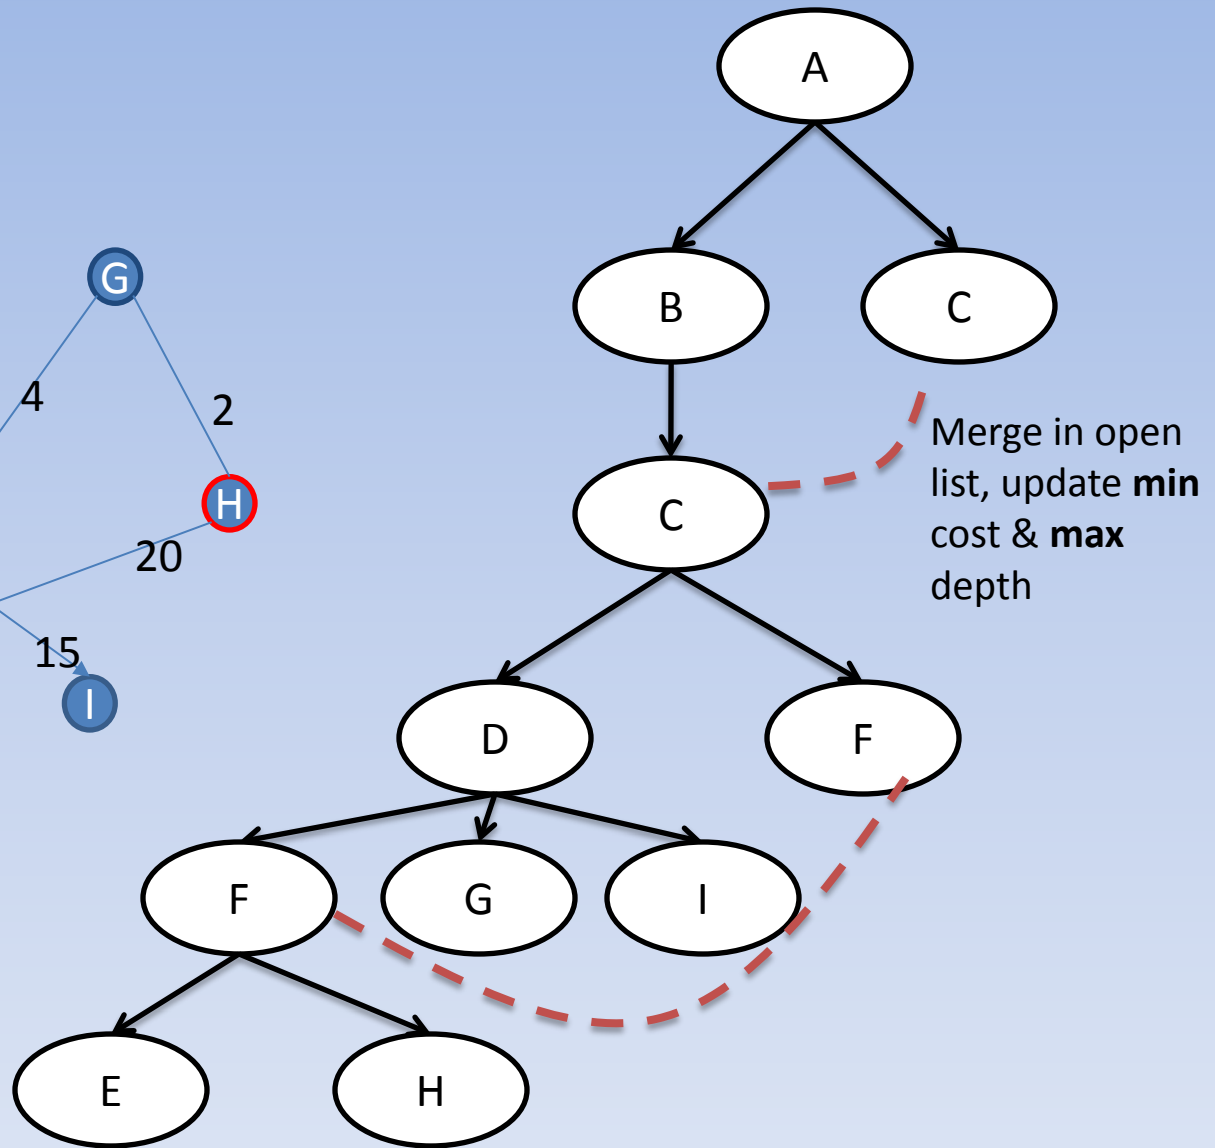

Uniform Cost Search

- Expand the node on the open list with the lowest path cost
 - otherwise same as BFS
- Can be implemented using a priority queue

Depth First Search

- Expands deepest nodes on open list first
- Can be implemented using a stack

Example



Iterative Deepening (ID-DFS)

- DFS may go down long, useless paths
 - We can parameterize it with a depth limit
 - If limit is reached, it will not expand further
- We can iteratively increase the depth parameter
 - Start with zero
 - If DFS with current depth fails, increment depth and restart

Bidirectional Search

- Why not search from both the initial state and the goal state?
- Idea: run simultaneous searches, checking for intersections between the open lists
 - If nonempty, path found
 - If using BFS at each end, and goal depth is d , time and space complexity will be reduced to $O(b^{d/2})$

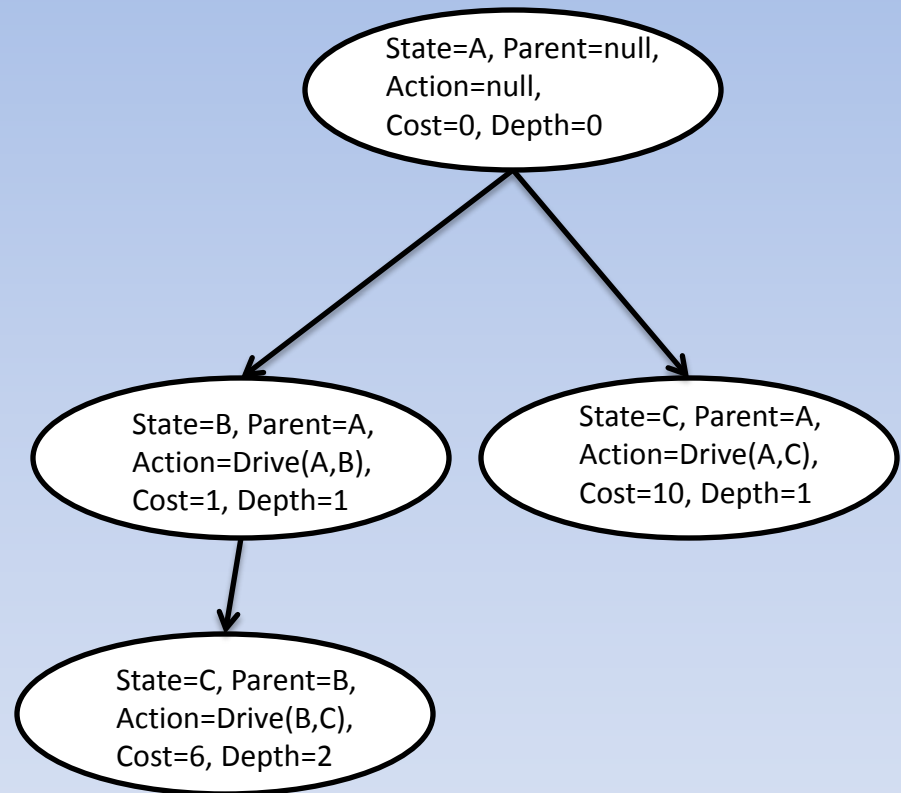

But...

- What if the goal is defined in terms of a predicate?
 - Could have lots of satisfying states, or could be hard to satisfy (SAT problem)
- Also constrains the operators
 - Note when searching from goal, need to find “predecessors” ---could be tricky!
 - “All states that led to this checkmate configuration”

Search Algorithm Analysis

- As it searches, the agent generates a “search tree”
- Each node in the search tree represents some state in the search space, with extra bookkeeping information
 - The parent node
 - The action that was applied at the parent
 - Path cost
 - Depth
- **NOTE:** The search tree is only meant to visualize the flow of computation. It does not correspond to a data structure you would maintain in practice.

Example: Search Tree



Different search algorithms can be compared using characteristics of the search tree they generate.

Characteristics of the Search Tree

- Branching factor b
 - Maximum number of successors of any node
- Goal depth d
 - Depth of the shallowest goal in the search tree
- Max Path Length m
 - Maximum length of a path in the search space

Algorithm Performance

- Completeness
 - Algorithm always finds a solution if it exists?
- Optimality
 - Algorithm always finds minimum cost solution?
- Time Complexity
 - Time taken to find a solution?
- Space Complexity
 - Memory required to find a solution?

BFS Performance

- Complete?
- Yes
- Optimal?
- Sometimes (when?)
- Time Complexity?
- $O(b^{d+1})$
- Space Complexity?
- $O(b^{d+1})$

DFS Performance

- Complete?
- Optimal?
- Time Complexity?
- Space Complexity?
- Yes, assuming no infinite paths
- Sometimes
- $O(b^m)$
- $O(b^m)$

Uninformed vs. Informed Search

- Previous search methods use various methods to pick which node to expand
- But which node would we *really* like to expand, path costs being equal?
 - The one that is *really the closest to the goal*
- But we don't know this
 - If we did, wouldn't need to search

“Informed” Search

- Uninformed or Blind Search
 - These algorithms only use the information in the problem setup to find a solution
- Informed or Heuristic Search
 - These algorithms also use an extra function, a *search heuristic*, to find a solution
 - The search heuristic is not part of the problem definition---it is up to the agent designer to specify it (some recent work on *learning* the heuristic)